


Robot Eval

A simple guide to using the platform — no technical background needed.



The Overview page — the first thing you see when you open Robot Eval.


Robot Eval is a website that helps teams understand how well their robots are doing. Robots practice tasks — picking things up, cooking, plugging in cables — and every practice attempt is recorded on video. This platform collects all those recordings in one place so you can watch them, see which attempts worked, spot what keeps going wrong, and check whether the robot is getting better over time.

Words you'll see (and what they really mean)

The platform uses a few robotics words. Here is a plain-English translation of each one:

Word on screen	What it means in plain English
Episode	One single attempt at a task, recorded on video. Like one 'take' in filming.
Dataset	A collection of episodes that belong together — e.g. all attempts at making coffee.
Task	What the robot was asked to do, e.g. 'Insert the cable into the port'.
Policy	The robot's 'brain' — the software that decides its movements. Policies have version numbers (like app updates), e.g. policy-v1.4.2.
human-teleop	A person remotely steered the robot for that episode (no AI brain involved). Teams record these to teach the robot by example.
Rollout	Another word for letting the robot attempt a task and recording it.
Success / Failure	Whether the attempt worked. Judged automatically or by a person.
Unscored	Nobody (and no software) has judged this attempt yet — the video exists, but it hasn't been marked as success or failure.
Failure category	The kind of mistake, e.g. 'Grasp slipped' (dropped its grip) or 'Collision' (bumped into something).
Success rate	Out of the attempts that were judged, the percentage that worked.
Benchmark pack	A themed group of tasks used to test robots, e.g. tabletop object handling.
Embodiment	Which physical robot was used — its model and body type.
DoF	'Degrees of freedom' — how many joints the robot can move. More joints, more dexterity. A human arm has 7.
Sensors / cameras	The robot's 'eyes' — e.g. a wrist camera or an overhead camera.
Coverage	How complete the recorded information is for an episode (100% = nothing missing).
Adapter	A translator that lets the platform read recordings from any robot software.
Source format	Which recording system the data originally came from (LeRobot, ROS 2, etc.).
Duration	How long the attempt took, in seconds.
Interventions	How many times a person had to step in and help mid-attempt.

Page 1 — Overview: the health dashboard



The Overview page summarises everything in one glance.

- **The four big numbers** at the top show: how many attempts are stored, what share of the judged ones succeeded, how long a typical attempt takes, and how often robots bump into things.
- **The red banner** appears automatically when a newer robot brain performs worse than the previous one — so problems are impossible to miss. Click 'Investigate in compare' to see details.
- **Rollouts by day** is a calendar chart: green = successful attempts, red = failed attempts, and the teal line tracks the success percentage day by day.
- **Failures by category** counts each kind of mistake. Clicking any row jumps straight to the matching videos.
- **Datasets** lists every collection of recordings loaded into the platform, including real public recordings from Stanford's Mobile ALOHA robot and others.
- **Recent episodes** shows the newest attempts. Click any ID (like ep_00053) to watch it.

Tip: almost everything on this page is clickable. If a number looks interesting, click it — you'll land on the exact list of attempts behind that number.

Page 2 — Episodes: browse and search every attempt



Robot Eval 0.1.0

robot eval & failure monitor

VIEWS

Overview

Episodes

Compare

Episodes

366 of 366 episodes

Task, instruction, or episode id...

Show every episode where:

Grasp slipped 4

Missed grasp 11

Dropped object 3

Wrong object 2

Collision 12

Stalled / timeout 5

Planning failure 1

OUTCOME

All Success Failure

DATASET

Tabletop pick-and-p...

22

Warehouse manipul...

42

Manual eval session...

32

PushT — T-block pu...

52

ALOHA static coffee...

50

SO-101 pick & place...

50

Mobile ALOHA — co...

18

Mobile ALOHA — wa...

50

RJ45 cable insertion...

50

SOURCE FORMAT

LeRobot

292

ROS 2 bag

42

CSV + video

32

POLICY VERSION

policy-v1.5.0-rc1

32

policy-v1.4.2

34

policy-v1.3.0

30

human-teleop

220

act-rj45-ckpt40k

50

BENCHMARK PACK

manipulation-table...

276

mobile-navigation

22

mobile-manipulation

68

The Episodes page — a searchable library of every recorded attempt.

- **Search box** (top right): type anything — a task like 'coffee', or an episode ID — and the list narrows as you type.
- **The grey buttons** under the heading ('Grasp slipped', 'Collision'...) are one-click shortcuts: press one to see every attempt that failed in that specific way. Press again to turn it off.
- **The left sidebar** has tick-boxes to narrow the list by dataset, recording system, robot-brain version, or task group. Tick several — they combine.
- **All / Success / Failure** buttons show everything, only wins, or only failures.
- **Clear all filters** resets everything with one click.
- Each row shows a small preview, the task, which robot, which brain version, the outcome, how long it took, and how complete the record is. **Click any row** to open the full episode.



Robot Eval 0.1.0

robot eval & failure monitor

VIEWS

Overview

Episodes

Compare

3 adapters online

self-hosted · mock data

schema v1.0

Episodes

12 of 366 episodes

Task, instruction, or episode id...

Show every episode where:

Grasp slipped 4

Missed grasp 11

Dropped object 3

Wrong object 2

Collision 12

Stalled / timeout 5

Planning failure 1

Clear all filters

OUTCOME

All

Success

Failure

DATASET

Tabletop pick-and-p... 22

Warehouse manipul... 42

Manual eval session... 32

PushT — T-block pu... 52

ALOHA static coffee... 50

SO-101 pick & place... 50

Mobile ALOHA — co... 18

Mobile ALOHA — wa... 50

RJ45 cable insertion... 50

SOURCE FORMAT

LeRobot 292

ROS 2 bag 42

CSV + video 32

POLICY VERSION

policy-v1.5.0-rc1 32

policy-v1.4.2 34

policy-v1.3.0 30

human-teleop 220

act-rj45-ckpt40k 50

BENCHMARK PACK

manipulation-table... 276

mobile-navigation 22

mobile-manipulation 68

EPISODE	TASK	EMBODIMENT	POLICY	OUTCOME	FAILURE	DURATION	COVERAGE
 ep_00005	Insert peg into hole manipulation-tabletop	UR5e	policy-v1.4.2	Failure	Collision	16.4s	71%
 ep_00025	Insert peg into hole manipulation-tabletop	SO-101	policy-v1.5.0-rc1	Failure	Collision	26.1s	80%
 ep_00027	Deliver item to marked zone mobile-navigation	Stretch RE2	policy-v1.3.0	Failure	Collision	8.3s	53%
 ep_00030	Open drawer, retrieve cup manipulation-tabletop	SO-101	policy-v1.4.2	Failure	Collision	16.5s	59%
 ep_00043	Navigate to charging dock mobile-navigation	Stretch RE2	policy-v1.4.2	Failure	Collision	40.9s	55%
 ep_00045	Insert peg into hole manipulation-tabletop	SO-101	policy-v1.3.0	Failure	Collision	9.1s	63%
 ep_00060	Insert peg into hole manipulation-tabletop	SO-101	policy-v1.5.0-rc1	Failure	Collision	43.7s	75%
 ep_00068	Deliver item to marked zone mobile-navigation	Stretch RE2	policy-v1.5.0-rc1	Failure	Collision	23.8s	53%
 ep_00073	Insert peg into hole manipulation-tabletop	SO-101	policy-v1.5.0-rc1	Failure	Collision	43.9s	87%
 ep_00074	Deliver item to marked zone mobile-navigation	Stretch RE2	policy-v1.5.0-rc1	Failure	Collision	30.5s	75%
 ep_00083	Navigate to charging dock mobile-navigation	Stretch RE2	policy-v1.4.2	Failure	Collision	33.5s	56%
 ep_00095	Navigate to charging dock mobile-navigation	Stretch RE2	policy-v1.4.2	Failure	Collision	43.1s	78%

After pressing the 'Collision' button: the list shrinks to only collision failures, and the count updates (here: 12 of 366).

Page 3 — Episode detail: watch one attempt up close



Robot Eval 0.1.0

robot eval & failure monitor

VIEWS

Overview

Episodes


Compare

Episodes / shrimp_ep_00000

Sauté and serve shrimp

"Sauté the raw shrimp on both sides, then serve it in the bowl."

Unscored human-teleop mobile-manipulation



observation.images.cam_high

real footage · 0s-75s of source file

observation.images.cam_high observation.images.cam_left_wrist observation.images.cam_right_wrist

View raw source

Action & state timeline

0.0s → 75.0s · synthetic signal preview

No anomaly detected

base_velocity

left_gripper_force

right_gripper_force

arm_sync_error

0.0s 75.0s

> CANONICAL EPISODE SCHEMA (JSON)

Outcome

success

determined by

Unscored

not recorded in source dataset

Metrics

duration

interventions

collisions

1m 15s

—

—

Embodiment

robot type

model

DoF

Mobile bimanual

Mobile ALOHA

14

Provenance

dataset

source format

schema version

recorded

coverage

Mobile ALOHA — cook shrimp (Stanford)

LeRobot


v1.0

Mar 10, 2026, 11:44 AM


80%

A real episode: Stanford's Mobile ALOHA robot cooking shrimp, streamed from the public dataset.

- **Video player:** press play to watch the robot's own camera footage of this exact attempt.
- **Grey labels under the video** list the robot's cameras for this recording.
- **Action & state timeline:** a wave-style picture of the robot's movements over time (marked 'synthetic preview' when real sensor traces aren't loaded yet).
- **Outcome** (right column): whether it worked, and importantly *how that was decided* — by software, by a person, or 'Unscored' if not judged yet.
- **Metrics:** duration, how often a person stepped in, and any collisions.
- **Embodiment:** which robot body did the work.
- **Provenance:** where this recording came from and how complete it is.



Page 4 — Compare: is the new robot brain better?



Compare puts two robot-brain versions side by side, like a before-and-after report.

- **Baseline and Candidate** drop-downs (top): pick the older version and the newer one you want to judge.
- **The four tiles** show the newer version's numbers, with green ▲ for improvements and red ▼ for regressions compared to the older version.
- **Failures by category** shows whether each kind of mistake got more or less common.
- **Regressions by task** is the heart of the page: tasks are sorted worst-first, so the biggest new problem is always the first row. 'View episodes' jumps to the actual videos so you can see the failures with your own eyes.

Everyday touches

- **Light / dark mode:** the sun/moon button at the bottom of the left sidebar switches the look. Your choice is remembered.
- **Shareable links:** the web address updates as you filter. Copy it, send it to a colleague, and they'll see exactly the same filtered view.
- **Real public data:** the platform streams real robot recordings directly from Hugging Face (a public library of robot data) — including Stanford's Mobile ALOHA cooking robot, an industrial cable-plugging arm, and a coffee-making robot.

That's everything. If you can browse an online video library and read a scoreboard, you can use Robot Eval.